

Generative Continuation Geometry

Working Draft

Flyxion

July 2026

Contents

1	Continuation Attractors	2
2	Continuation Quotients	6
3	Distinction Loss Under Quotients	11
4	Continuation Operators	14
5	Toward Universal Generation	17
	What Remains	20
	References	22

Chapter 1

Continuation Attractors

A fern frond is not drawn leaf by leaf. Look closely at one pinna and it is, near enough, a smaller fern; look closer still at one pinnule of that pinna and it is, near enough, a smaller fern again. Nobody specified the shape of the whole plant directly. A small number of growth rules, applied to their own output over and over, produced something with more detail than any of the rules themselves contain. The shape is not stored anywhere. It is generated.

Continuation Geometry spent its length measuring futures: how much room a state has, how far apart two states' futures sit, when a distinction between futures collapses to nothing. All of that took the space of futures as given, something to be measured once it already existed. This book asks a different question. Can a space of futures itself be grown, the way the fern is grown, from a small number of simple moves applied to their own output? If continuation spaces can be generated rather than merely measured, the natural unit of study stops being the space and becomes the small set of operations that built it.

This opening chapter answers that question in the one case where the answer is already, in a precise sense, classical. It is not a new theorem. It is an old and beautiful one, due to Hutchinson, applied here to the specific metric geometry the first book already built. The payoff is that self-similar continuation structures, the kind suggested by a duck, a crown, and a leaf all arising from the same underlying rule, turn out to rest on exactly the Hausdorff-distance machinery already in hand, with nothing new invented and nothing smuggled in from outside.

Preliminaries: expansion is not generation

Two facts from *Continuation Geometry* are worth restating here, briefly and as what they actually are: lemmas, not theorems, and useful precisely because they show what generation is *not*.

Lemma 1.1 (Reachable expansion). *For $S_t \subseteq X$ and $\text{Adj}(S_t) := \bigcup_{x \in S_t} \Phi(x)$ under a one-step reachability map Φ , the sequence $S_{t+1} := S_t \cup \text{Adj}(S_t)$ satisfies $S_t \subseteq S_{t+1}$ for every t .*

Proof. Immediate: $S_t \subseteq S_t \cup \text{Adj}(S_t) = S_{t+1}$ by definition of union. □

Remark 1.1. *This, together with the reachable-set recursion proved by induction in *Continuation Geometry* (Kauffman's adjacent possible as an instance of finite-horizon reachability), is the full content of what that book called generation, and it is worth being honest that it earns the name only loosely. The*

sequence $S_0, S_1, S_2, \dots$ grows monotonically and without bound unless the underlying space is finite; nothing about it converges, and nothing about it is self-similar. It is accumulation, not generation in the sense the fern illustrates. What follows is different in kind: rather than a sequence of ever-larger unions, it is a fixed operator applied repeatedly to itself, and the interesting behavior is not growth but convergence to a shape that reproduces its own generating rule.

The space of shapes

Fix $X = \mathbb{R}^n$ (or any complete metric space) and let $\mathcal{K}(X)$ denote the set of nonempty compact subsets of X , equipped with the Hausdorff distance d_H of Chapter 6 of *Continuation Geometry*.

Theorem 1.1 (Completeness of the space of shapes). *If (X, d) is complete, then $(\mathcal{K}(X), d_H)$ is complete.*

Proof sketch. Given a Cauchy sequence $\{A_n\}$ in $(\mathcal{K}(X), d_H)$, define $A := \bigcap_{N \geq 1} \overline{\bigcup_{n \geq N} A_n}$. One shows, using completeness of X and the Cauchy property of $\{A_n\}$ in d_H , that A is nonempty and compact, and that $d_H(A_n, A) \rightarrow 0$. This is a standard but genuinely technical result in its full form, requiring care with total boundedness and a diagonal argument to extract convergent selections from each A_n ; the complete argument is given in [1] and [2] and is not reproduced in full here. $\square$

Remark 1.2. *This theorem is doing real work that is easy to take for granted. Without it, nothing built on top of $(\mathcal{K}(X), d_H)$ would be guaranteed to have a limit at all; a sequence of shapes could get closer and closer together, in the sense of Hausdorff distance, without converging to any actual shape. That this cannot happen, that the space of compact reachable sets is itself complete whenever the underlying state space is, is the entire reason the construction below is guaranteed to produce something rather than merely getting closer to nothing in particular.*

Contractions and the Hutchinson operator

Definition 1.1 (Contraction). *$T : X \rightarrow X$ is a contraction with constant $c \in [0, 1)$ if $d(Tx, Ty) \leq c d(x, y)$ for all $x, y \in X$.*

Lemma 1.2 (Contractions act contractively on shapes). *If T is a contraction on X with constant c , then for $A, B \in \mathcal{K}(X)$, $d_H(T(A), T(B)) \leq c d_H(A, B)$, where $T(A) := \{Ta : a \in A\}$.*

Proof. Fix $a \in A$. Since B is compact, there is $b \in B$ with $d(a, b) = \text{dist}(a, B) \leq d_H(A, B)$. Then $d(Ta, Tb) \leq c d(a, b) \leq c d_H(A, B)$, and since $Tb \in T(B)$, $\text{dist}(Ta, T(B)) \leq c d_H(A, B)$. Taking the supremum over $a \in A$ gives $\sup_{a' \in T(A)} \text{dist}(a', T(B)) \leq c d_H(A, B)$. The symmetric argument, exchanging the roles of A and B , gives the same bound the other way, so $d_H(T(A), T(B)) \leq c d_H(A, B)$. $\square$

Lemma 1.3 (Unions are 1-Lipschitz). *For finite collections $\{A_i\}_{i=1}^n, \{B_i\}_{i=1}^n$ of nonempty compact sets, $d_H(\bigcup_i A_i, \bigcup_i B_i) \leq \max_i d_H(A_i, B_i)$.*

Proof. Let $m := \max_i d_H(A_i, B_i)$. For $x \in \bigcup_i A_i$, $x \in A_j$ for some j , so $\text{dist}(x, B_j) \leq d_H(A_j, B_j) \leq m$, and since $B_j \subseteq \bigcup_i B_i$, $\text{dist}(x, \bigcup_i B_i) \leq \text{dist}(x, B_j) \leq m$. Taking the supremum over $x \in \bigcup_i A_i$ gives $\sup_x \text{dist}(x, \bigcup_i B_i) \leq m$. The symmetric argument, exchanging the roles of the two collections, gives the same bound the other way, so $d_H(\bigcup_i A_i, \bigcup_i B_i) \leq m$. $\square$

Definition 1.2 (Hutchinson operator). Given contractions $T_1, \dots, T_n$ on X with constants $c_1, \dots, c_n$, define $\mathcal{H} : \mathcal{K}(X) \rightarrow \mathcal{K}(X)$ by $\mathcal{H}(A) := \bigcup_{i=1}^n T_i(A)$.

Proposition 1.1 ($\mathcal{H}$ is a contraction on shapes). $d_H(\mathcal{H}(A), \mathcal{H}(B)) \leq c d_H(A, B)$ for $c := \max_i c_i < 1$.

Proof. By Lemma 1.3, $d_H(\mathcal{H}(A), \mathcal{H}(B)) = d_H(\bigcup_i T_i(A), \bigcup_i T_i(B)) \leq \max_i d_H(T_i(A), T_i(B))$. By Lemma 1.2, $d_H(T_i(A), T_i(B)) \leq c_i d_H(A, B)$ for each i . Hence $d_H(\mathcal{H}(A), \mathcal{H}(B)) \leq \max_i c_i d_H(A, B) = c d_H(A, B)$. $\square$

The Continuation Attractor Theorem

Theorem 1.2 (Continuation Attractor Theorem, after Hutchinson). Let $T_1, \dots, T_n$ be contractions on a complete metric space X with constants $c_1, \dots, c_n$, and let $c := \max_i c_i < 1$. Then there exists a unique $A^* \in \mathcal{K}(X)$ with $\mathcal{H}(A^*) = A^*$, and for any $A_0 \in \mathcal{K}(X)$,

$$d_H(\mathcal{H}^k(A_0), A^*) \leq \frac{c^k}{1-c} d_H(A_0, \mathcal{H}(A_0)) \xrightarrow{k \rightarrow \infty} 0.$$

Proof. By Theorem 1.1, $(\mathcal{K}(X), d_H)$ is complete. By Proposition 1.1, $\mathcal{H}$ is a contraction on this complete metric space. The Banach Fixed Point Theorem then supplies a unique fixed point A^* and the stated geometric convergence rate, applied directly to $\mathcal{H}$ on $(\mathcal{K}(X), d_H)$. $\square$

Remark 1.3. No part of this proof required inventing anything: completeness of $\mathcal{K}(X)$ is Theorem 1.1, the contraction property of $\mathcal{H}$ is Proposition 1.1 built from two short lemmas, and the conclusion is the Banach Fixed Point Theorem, unmodified. This is the sense in which this chapter is mathematically conservative even though its subject, self-similar generated shape, sounds exotic: every piece was already sitting in the metric structure Chapter 6 of Continuation Geometry built for an entirely different purpose.

A worked attractor

Example 1.1 (Two refusals). Let $X = [0, 1]$, and let $T_1(x) := x/3$, $T_2(x) := x/3 + 2/3$, both contractions with constant $c = 1/3$. The attractor A^* satisfying $A^* = T_1(A^*) \cup T_2(A^*)$ is the standard middle-thirds Cantor set.

Remark 1.4. Read in the vocabulary of continuation, T_1 and T_2 are the two survivors of a repeated refusal: at every stage, whatever occupies the middle third of the current set of admissible outcomes is discarded, and the remaining two-thirds are each rescaled to stand in for the whole. What survives in the limit is neither of the two pieces nor their simple union at any finite stage, but the unique shape that reproduces itself exactly under the same rule that built it, uncountable, nowhere dense, and entirely determined by two contraction maps. Whether operators named Collapse or Refuse elsewhere in this corpus are literally instances of T_i in this technical sense, maps satisfying the contraction inequality on some genuine underlying metric, is a claim about those specific constructions and is not asserted here; what is established is only that if a family of continuation-narrowing operators can be shown to be contractions on $(\mathcal{K}(X), d_H)$, this theorem applies to them without modification.

What this chapter does not yet claim

This chapter establishes that self-similar continuation structures are mathematically available, not that any particular structure this corpus has described informally is an instance of one. Two gaps in particular are worth naming rather than leaving implicit. First, nothing here says which operators arising from an actual continuation problem, as opposed to a hand-picked pair of affine maps on $[0, 1]$, are genuinely contractions on a genuine metric; that has to be checked case by case, the way it would for any application of a fixed-point theorem. Second, and more speculatively, the question of whether a small, fixed family of primitive operators could generate *every* admissible continuation geometry to arbitrary Hausdorff-distance accuracy, a continuation-geometric analogue of universal approximation, is not addressed by this theorem and is not claimed here. It is a genuine open question, and answering it is a natural place for this book to go next, once the machinery in the chapters between here and there, quotient maps and operator composition among them, has been built on equally solid ground.

Chapter 2

Continuation Quotients

Two rivers can rise on opposite sides of a mountain range, cross entirely different terrain, pass through valleys that share nothing in common, and still arrive at the same confluence and become, from that point on, one river. Everything about their separate histories, which rocks each one wore smooth, which tributaries fed it, how long each took to descend, becomes irrelevant downstream of the merge. A boat launched below the confluence cannot tell which branch its water came from. Not because the information was hidden. Because it no longer makes a difference to anything the boat will experience from here on.

The previous chapter showed that a continuation space can be grown from a small number of generating operators. The natural next question is what happens when two different generating histories, different sequences of operators, different starting points, produce the same continuation space at the end. When does that happen, and what, if anything, does it mean? The temptation is to reach immediately for language like information loss or memory collapse, language already used loosely elsewhere in this corpus. This chapter resists that temptation on purpose. It defines the relevant object first, proves what can be proved about it without reference to any interpretation, and only afterward asks what, if anything, has been lost.

The continuation quotient

Definition 2.1 (Continuation quotient). *Let H be a set of histories and let $\Gamma : H \rightarrow \mathcal{K}(X)$ assign to each history its continuation space (or, following Chapter 6 of Continuation Geometry, its finite-horizon reachable set R_T for a fixed horizon T). Define $h_1 \sim_\Gamma h_2$ iff $\Gamma(h_1) = \Gamma(h_2)$. This is an equivalence relation (reflexivity, symmetry, and transitivity all following immediately from equality of sets), and the continuation quotient is $H/\sim_\Gamma$, with quotient map $q : H \rightarrow H/\sim_\Gamma$ sending h to its equivalence class $[h]$.*

Nothing has been said yet about distinctions, memory, or loss. This is deliberate: the object exists, and has structure, independently of any of those interpretations.

Proposition 2.1 (Well-defined continuation projection). *There exists a unique map $\bar{\Gamma} : H/\sim_\Gamma \rightarrow \mathcal{K}(X)$ such that $\Gamma = \bar{\Gamma} \circ q$.*

Proof. Define $\bar{\Gamma}([h]) := \Gamma(h)$. This is well-defined: if $[h_1] = [h_2]$, then $h_1 \sim_\Gamma h_2$, so $\Gamma(h_1) = \Gamma(h_2)$ by definition of $\sim_\Gamma$, so the value assigned to $[h_1] = [h_2]$ does not depend on which representative was chosen. By construction, $\bar{\Gamma}(q(h)) = \bar{\Gamma}([h]) = \Gamma(h)$ for every $h \in H$, so

$\Gamma = \bar{\Gamma} \circ q$. For uniqueness, suppose $\bar{\Gamma}'$ also satisfies $\Gamma = \bar{\Gamma}' \circ q$. Since q is surjective, every element of $H/\sim_\Gamma$ is $q(h)$ for some h , and $\bar{\Gamma}'(q(h)) = \Gamma(h) = \bar{\Gamma}(q(h))$ for every such h , so $\bar{\Gamma}' = \bar{\Gamma}$. $\square$

Remark 2.1. *This is routine, and it is worth saying plainly that it is routine: it is the universal property of a quotient set, applied to nothing more exotic than the map Γ itself. What makes it worth stating is not its difficulty but what it licenses: from here on, $\bar{\Gamma}$ is a genuine function on the quotient, not merely a suggestive notation, and anything proved about $\bar{\Gamma}$ is a fact about the quotient set itself, not about any particular history used to reach a point in it.*

A structural lemma

The worked example below relies on the following fact about time-homogeneous, deterministic-image dynamics, worth isolating here.

Lemma 2.1 (One-step agreement propagates). *Suppose $R_1(x) = R_1(y)$ for states x, y under a time-homogeneous reachability map. Then $R_T(x) = R_T(y)$ for every $T \geq 1$.*

Proof. By induction on T . The case $T = 1$ is the hypothesis. Suppose $R_T(x) = R_T(y)$. Since the dynamics are time-homogeneous, $R_{T+1}(x) = \bigcup_{z \in R_T(x)} R_1(z)$ and likewise for y ; since $R_T(x) = R_T(y)$, these unions range over the same set of states z , so $R_{T+1}(x) = R_{T+1}(y)$. $\square$

Remark 2.2. *This is the precise sense in which a collision, once it happens, is permanent under time-homogeneous dynamics: two states whose immediate futures coincide exactly will have every subsequent future coincide as well, since there is nothing left in either state's future to distinguish it from the other's.*

Quotient invariants

Proposition 2.2 (Quotient invariance). *Let $F : \mathcal{K}(X) \rightarrow Y$ be any function. Then $F \circ \Gamma : H \rightarrow Y$ factors through q : there is a unique $\bar{F} : H/\sim_\Gamma \rightarrow Y$ with $F \circ \Gamma = \bar{F} \circ q$.*

Proof. $F \circ \Gamma : H \rightarrow Y$ satisfies $(F \circ \Gamma)(h_1) = (F \circ \Gamma)(h_2)$ whenever $h_1 \sim_\Gamma h_2$, since $\Gamma(h_1) = \Gamma(h_2)$ by definition of $\sim_\Gamma$ and F is a function, hence single-valued. Apply Proposition 2.1 with $F \circ \Gamma$ in place of Γ to obtain the factorization. $\square$

Corollary 2.1. *Continuation volume Vol_T , the distinction margin A_D^T relative to any fixed reference distinction, and boundary proximity δ , all defined in Continuation Geometry as functions of a reachable set or of a state's position relative to one, descend to well-defined functions on $H/\sim_\Gamma$.*

Proof. Each is of the form $F(\Gamma(h))$ or $F(R_T(x(h)))$ for a function F depending only on the reachable set itself (Lebesgue measure of the set, Hausdorff distance to another fixed set, ambient distance from a fixed reference set), so Proposition 2.2 applies directly to each. $\square$

Remark 2.3. *This is a broader statement than it might first appear. It says that any quantity this book or its predecessor has defined, or will define, purely in terms of a reachable set, without reference to which history produced that set, is automatically constant on $\sim_\Gamma$ -equivalence classes, with no further argument required each time a new such quantity is introduced. The work of checking invariance, done once here, does not need to be redone for every future functional built the same way.*

Finite checkability

Lemma 2.1 shows that equal one-step images propagate forward forever. It says nothing about the converse: two states whose one-step images differ can still coincide at some later horizon. Nothing in this chapter has claimed otherwise, and the general question, whether continuation equivalence at every horizon can still be checked in finitely many steps despite this, deserves a correct answer rather than an implied one.

Proposition 2.3 (Eventual periodicity bounds the check). *If X is finite, then for any $x \in X$ the sequence $R_1(x), R_2(x), R_3(x), \dots$ is eventually periodic, with both the pre-period and the period bounded by $2^{|X|}$. Consequently, whether $R_T(x) = R_T(y)$ holds for all sufficiently large T can be determined by comparing $R_T(x)$ and $R_T(y)$ for T ranging only up to $2 \cdot 2^{|X|}$.*

Proof. Each $R_T(x)$ is an element of 2^X , which has $2^{|X|}$ elements, so the sequence $R_1(x), R_2(x), \dots$ takes at most $2^{|X|}$ distinct values; among its first $2^{|X|} + 1$ terms, two must coincide by pigeon-hole, $R_p(x) = R_q(x)$ for some $1 \leq p < q \leq 2^{|X|} + 1$. Time-homogeneity means $R_{T+1} = \Psi(R_T)$ for the fixed map $\Psi(S) := \bigcup_{z \in S} \Phi(z)$, so $R_p(x) = R_q(x)$ gives $R_{p+j}(x) = \Psi^j(R_p(x)) = \Psi^j(R_q(x)) = R_{q+j}(x)$ for every $j \geq 0$: the sequence is periodic with period dividing $q - p \leq 2^{|X|}$ from index $p \leq 2^{|X|}$ onward. The same bound applies to the sequence generated by y . Once both sequences have entered their eventual cycles, comparing them over one further full period, at most $2^{|X|}$ more steps, settles whether they agree from that point on. $\square$

Remark 2.4. *This is the correct version of the claim that continuation equivalence is finitely checkable. It is a coarser bound than a sharp analysis of any particular system would give, since $2^{|X|}$ is a worst case over all possible dynamics on a state space of that size, but it is unconditional: no structural assumption on Φ beyond finiteness of X and time-homogeneity is needed. The traffic light below reaches its own eventual behavior in four steps, far inside this bound, and it is worth seeing exactly how.*

The traffic light, revisited

Recall the traffic light of *Continuation Geometry*, Chapter 10: states $(k, r) \in \{0, 1, 2, 3\} \times \{0, 1\}$, phase k and pending request r . To compute Γ properly requires treating the system's branching explicitly rather than leaving it implicit. Let $e \in \{0, 1\}$ record whether a new pedestrian request arrives during a given step (exogenous, not controlled by the phase). The transition is

$$\Phi((k, r), e) = ((k + 1) \bmod 4, r'), \quad r' = \begin{cases} 0 & k + 1 \equiv 0 \pmod{4} \text{ (clear on arrival at red)} \\ 1 & k + 1 \not\equiv 0 \pmod{4} \text{ and } (r = 1 \text{ or } e = 1) \\ 0 & k + 1 \not\equiv 0 \pmod{4} \text{ and } r = 0 \text{ and } e = 0. \end{cases}$$

Let H be the set of finite paths of the system, and $\Gamma(h) := R_1(x(h))$, the one-step reachable set from the endpoint of h , ranging over $e \in \{0, 1\}$.

Proposition 2.4 (Exact computation of the traffic-light quotient at $T = 1$). *Under $\sim_\Gamma$, the eight states fall into exactly seven classes: the single nontrivial class $\{(3, 0), (3, 1)\}$, and six singleton classes, one for each remaining state.*

Proof. Compute R_1 for all eight states directly from Φ . For $k \neq 3$, arrival is at a nonzero phase and the ordinary branching rule applies, $r' = 1$ iff $r = 1$ or $e = 1$:

$$\begin{aligned} R_1(0, 0) &= \{(1, 0), (1, 1)\}, & R_1(0, 1) &= \{(1, 1)\}, \\ R_1(1, 0) &= \{(2, 0), (2, 1)\}, & R_1(1, 1) &= \{(2, 1)\}, \\ R_1(2, 0) &= \{(3, 0), (3, 1)\}, & R_1(2, 1) &= \{(3, 1)\}. \end{aligned}$$

For $k = 3$, arrival is at phase 0, so clearing fires unconditionally regardless of r or e :

$$R_1(3, 0) = \{(0, 0)\}, \quad R_1(3, 1) = \{(0, 0)\}.$$

Comparing all eight sets directly: the three size-two sets, $\{(1, 0), (1, 1)\}$, $\{(2, 0), (2, 1)\}$, $\{(3, 0), (3, 1)\}$, are pairwise distinct; the four remaining size-one sets, $\{(1, 1)\}$, $\{(2, 1)\}$, $\{(0, 0)\}$ (appearing twice, from $(3, 0)$ and $(3, 1)$), and $\{(0, 0)\}$ again, give exactly one coincidence, $R_1(3, 0) = R_1(3, 1) = \{(0, 0)\}$, and no other. $\square$

Now define the mechanical quotient $\sim_M$ independently of Γ , as the identification the design actually performs: the clearing rule fires on the transition *into* phase 0, so it is the two states about to make that transition, $(3, 0)$ and $(3, 1)$, that the design treats as interchangeable, both being followed unconditionally by $(0, 0)$.

Proposition 2.5 (The design is exact at $T = 1$). $\sim_M = \sim_\Gamma$.

Proof. Both relations have exactly one nontrivial class, $\{(3, 0), (3, 1)\}$, and agree as singletons on every other state: $\sim_M$ by definition of the clearing rule, $\sim_\Gamma$ by Proposition 2.4. Two equivalence relations on the same finite set with identical partitions are equal. $\square$

Remark 2.5. *It is worth being precise about which of two similar-sounding designs this result actually vindicates. Clearing on arrival at phase 0 (the rule above) is exact. A superficially similar alternative, clearing whenever the system is currently at phase 0 rather than arriving there, would instead identify $(0, 0)$ and $(0, 1)$. But $R_1(0, 0) = \{(1, 0), (1, 1)\} \neq \{(1, 1)\} = R_1(0, 1)$: a state at the red light with no pending request can still branch toward either outcome next, while one with a pending request cannot, precisely because that request is about to be served. The alternative design would merge two states with genuinely different one-step futures, exactly the coarser, information-destroying case this chapter opened by distinguishing from the exact one. The two designs differ only in whether clearing is keyed to the transition or to the state, and only one of them is correct.*

What one horizon hides

Restricting Proposition 2.4 to $T = 1$ was not a simplification; it is what Γ was defined to measure. But Proposition 2.3 guarantees the longer-horizon question has a definite, finitely checkable answer, and for this system that answer is worth computing directly rather than left as an abstract guarantee.

Proposition 2.6 (The request bit is eventually redundant everywhere). *For each phase k , the pair $(k, 0)$ and $(k, 1)$ satisfies $R_T(k, 0) = R_T(k, 1)$ for all $T \geq 4 - k$, and not before.*

Proof by direct computation. By Proposition 2.4, $(3, 0) \sim_{\Gamma} (3, 1)$ at $T = 1 = 4 - 3$. Applying Ψ once more, $R_2(2, 0) = \Psi(R_1(2, 0)) = \Psi(\{(3, 0), (3, 1)\})$ and $R_2(2, 1) = \Psi(R_1(2, 1)) = \Psi(\{(3, 1)\})$; since $\Phi(3, 0, e) = \Phi(3, 1, e) = (0, 0)$ for both e , both unions reduce to $\{(0, 0)\}$, so $R_2(2, 0) = R_2(2, 1) = \{(0, 0)\}$, matching $T = 2 = 4 - 2$. The same one-step displacement argument gives $R_3(1, 0) = R_3(1, 1)$ at $T = 3 = 4 - 1$ and $R_4(0, 0) = R_4(0, 1)$ at $T = 4 = 4 - 0$: each pair's images are, one step earlier, exactly the pair merged at the previous phase, so equality propagates forward by exactly one horizon per phase step. That equality does not hold earlier follows from the cardinality mismatch already established at $T = 1$ in Proposition 2.4, propagated the same one step at a time: at $T = 4 - k - 1$, the images of $(k, 0)$ and $(k, 1)$ are the not-yet-merged pair one phase ahead. $\square$

Remark 2.6. *This is the sharper and more interesting fact the single-horizon result was hiding. The request bit is not redundant in general; it is redundant by a deadline, and that deadline is exactly the number of steps remaining until the light is guaranteed to cycle through its clearing transition, four steps for a state that has just left it, one step for a state about to enter it. Every state's request bit stops mattering to its continuation space at precisely the moment continuation stops being able to tell whether that request will still be pending when clearing finally arrives, and not one step before. The design clears the bit at the earliest possible moment, $T = 1$ after arrival, which is also the latest moment at which failing to clear it would still have been observably wrong; every other phase's bit becomes redundant on its own, without any design intervention, simply because the light's own periodicity eventually answers the question the bit was tracking.*

What quotients do not yet explain

This chapter has shown when two histories are interchangeable from the standpoint of every quantity built purely from continuation spaces, and it has shown, in one worked case, that a designed identification can be checked against that standard exactly rather than merely defended by intuition. It has not yet said anything about what is lost when a quotient is coarser than $\sim_{\Gamma}$, only that the traffic light's particular quotient happens not to be. The next chapter takes up the coarser case directly: not whether information is destroyed, which this chapter's machinery can already detect, but how much, and by what measure.

Chapter 3

Distinction Loss Under Quotients

A library catalog entry is a quotient of a book: title, author, and a handful of subject headings standing in for everything between the covers. A good catalog loses nothing a patron actually needs in order to decide whether to pull the book from the shelf. A bad one can lose quite a lot; a catalog that records only “fiction” or “nonfiction” will happily file a memoir next to a novel and a field guide next to a textbook, and a patron relying on that heading alone has no way to tell them apart until the book is already in hand. The catalog entry is not wrong, exactly; it simply was not built to preserve the distinctions this particular patron cares about.

The previous chapter showed how to check whether a given identification of histories is exact: whether it merges only histories whose continuation spaces already coincide, or whether it merges some that do not. That check was binary, exact or not. This chapter asks the sharper question a librarian actually has to answer. When an identification is not exact, how much is lost, and where exactly does the loss occur? The traffic light’s badly designed alternative from the previous chapter turns out to have a precise answer, not just a qualitative one.

Fibers and their content

Definition 3.1 (Fiber, Γ -content, loss). *For a quotient $q : H \rightarrow O$ and $o \in O$, the fiber over o is $q^{-1}(o) \subseteq H$. Its Γ -content is $C_\Gamma(o) := \{\Gamma(h) : h \in q^{-1}(o)\}$, the set of distinct continuation spaces appearing in the fiber. The loss at o is*

$$\mathcal{L}(o) := \log |C_\Gamma(o)|,$$

so $\mathcal{L}(o) = 0$ exactly when every history mapping to o under q shares a single continuation space.

Proposition 3.1 (Zero loss characterizes Γ -respecting quotients). $\mathcal{L}(o) = 0$ for every $o \in O$ if and only if q refines $\sim_\Gamma$: $h_1 \sim_q h_2 \implies h_1 \sim_\Gamma h_2$.

Proof. If q refines $\sim_\Gamma$, then $h_1 \sim_q h_2$ implies $\Gamma(h_1) = \Gamma(h_2)$, so every fiber of q lies inside a single $\sim_\Gamma$ -class and has $|C_\Gamma(o)| = 1$, giving $\mathcal{L}(o) = 0$. Conversely, if $\mathcal{L}(o) = 0$ for every o , then every fiber has $|C_\Gamma(o)| = 1$, so any two histories in the same fiber share the same Γ -value: $h_1 \sim_q h_2 \implies \Gamma(h_1) = \Gamma(h_2) \implies h_1 \sim_\Gamma h_2$, so q refines $\sim_\Gamma$. $\square$

The coarsest lossless quotient

Theorem 3.1 ($\sim_\Gamma$ is the coarsest zero-loss quotient). *Every zero-loss quotient $q : H \rightarrow O$ refines $\sim_\Gamma$, and $\sim_\Gamma$ itself has zero loss everywhere. Consequently $\sim_\Gamma$ is, up to relabeling of O , the unique coarsest quotient of H with zero loss.*

Proof. The first claim is Proposition 3.1. That $\sim_\Gamma$ itself has zero loss is immediate: by definition, $h_1 \sim_\Gamma h_2$ already means $\Gamma(h_1) = \Gamma(h_2)$, so every $\sim_\Gamma$ -class has exactly one Γ -value. Since every zero-loss q refines $\sim_\Gamma$, no zero-loss quotient can be strictly coarser than $\sim_\Gamma$, and $\sim_\Gamma$ achieves zero loss, so it is the coarsest one. $\square$

Remark 3.1. *This is the quotient-theoretic form of a fact Chapter 5 of Continuation Geometry established for sufficient statistics without this vocabulary. A sufficient statistic φ there was required to satisfy $\varphi(H_t) = \varphi(H'_t) \implies \Gamma(x_t | H_t) = \Gamma(x_t | H'_t)$, which is exactly the condition that the quotient induced by φ refines $\sim_\Gamma$. Theorem 3.1 says $\sim_\Gamma$ itself is the coarsest possible sufficient statistic: the most compressed summary of a history that still loses nothing continuation-relevant. Any coarser summary is not merely less detailed; by Proposition 3.1, it has strictly positive loss somewhere.*

Proposition 3.2 (Coarsening never decreases loss). *If q' refines q (every q' -fiber is contained in some q -fiber), then for every q -fiber $q^{-1}(o)$, expressed as a union of q' -fibers $q'^{-1}(o'_1), \dots, q'^{-1}(o'_k)$,*

$$\mathcal{L}_q(o) \geq \max_i \mathcal{L}_{q'}(o'_i).$$

Proof. $C_\Gamma(o) = \bigcup_i C_\Gamma(o'_i)$, since the Γ -values appearing in $q^{-1}(o)$ are exactly those appearing in its constituent q' -fibers. A union of sets is at least as large as any one of them, so $|C_\Gamma(o)| \geq |C_\Gamma(o'_i)|$ for each i , and log is monotone increasing. $\square$

Remark 3.2. *Merging two fibers can only add to the set of continuation spaces an observer sees when looking at the merged fiber; it can never remove one. This is the precise sense in which coarsening a quotient is a one-way trip toward loss: refining a lossy quotient can repair the loss, by splitting a mixed fiber back into pieces that each see only one Γ -value, but no further coarsening of an already-lossy quotient can undo the damage, only add to it or leave it unchanged.*

Distinction margin as a local loss detector

Proposition 3.3 (Vanishing margin is fiber collision). *For a distinction $\mathcal{D}$ and states x, y in different cells of $\mathcal{D}$, $A_{\mathcal{D}}^T(x) = 0$ (Chapter 7 of Continuation Geometry) if and only if x and y lie in the same fiber of the quotient induced by $\mathcal{D}$ but in different $\sim_\Gamma$ -classes at horizon T : that is, $\mathcal{D}$ merges x and y while $R_T(x) \neq R_T(y)$ for some witnessing y , or $R_T(x) = R_T(y)$ for the specific y achieving the infimum.*

Proof. Recall $A_{\mathcal{D}}^T(x) := \inf_{y \notin D(x)} d_H(R_T(x), R_T(y))$. This infimum is 0 exactly when some y outside x 's $\mathcal{D}$ -cell has $R_T(x) = R_T(y)$ (attained, under the compactness hypothesis already used in Chapter 7), which is exactly the statement that x and that y share a Γ -value at horizon T despite $\mathcal{D}$ having placed them in different cells: they are $\sim_\Gamma$ -equivalent (at horizon T) even though $\mathcal{D}$ distinguishes them. $\square$

Remark 3.3. *This gives the distinction margin of the first book a second reading. It was introduced there as a measure of how close a state sits to losing a distinction. Read through this chapter's vocabulary, it is something more specific: a pointwise diagnostic for whether the observer's chosen distinction $\mathcal{D}$, treated as its own quotient of the state space, disagrees with $\sim_\Gamma$ at that particular point. $A_{\mathcal{D}}^T(x) = 0$ does not merely mean x is close to indistinguishable from something outside its cell; by this chapter's Proposition 3.1 applied locally, it means $\mathcal{D}$'s quotient has positive loss exactly there, whether or not it has loss anywhere else.*

The traffic light, quantified

Return to the two designs compared in the previous chapter: the correct one, clearing on arrival at phase 0, and the alternative, clearing whenever currently at phase 0.

Example 3.1 (Zero loss, and one bit). *The correct design's quotient is $\sim_\Gamma$ itself (Proposition 4 of the previous chapter), so by Theorem 3.1 its loss is $\mathcal{L}(o) = 0$ at every fiber o . The alternative design merges $(0, 0)$ and $(0, 1)$ into a single fiber o_{bad} . Their Γ -values at $T = 1$ are $R_1(0, 0) = \{(1, 0), (1, 1)\}$ and $R_1(0, 1) = \{(1, 1)\}$, which are distinct, so $C_\Gamma(o_{\text{bad}}) = \{\{(1, 0), (1, 1)\}, \{(1, 1)\}\}$, a set of size 2, giving*

$$\mathcal{L}(o_{\text{bad}}) = \log 2,$$

exactly one bit, and $\mathcal{L}(o) = 0$ at every other fiber, since the alternative design agrees with the correct one everywhere except at this single merge.

Remark 3.4. *This is what the previous chapter's qualitative verdict, that the alternative design was coarser and information-destroying, actually amounts to once measured: not a diffuse degradation spread across the whole system, but a single, exactly located bit of loss, sitting at precisely the fiber where the alternative design merged two states that a patron of this particular library, a controller trying to plan four steps ahead, would still have wanted to tell apart. Every other fiber of the alternative design is, by coincidence or by the limited scope of its one design error, exactly as good as the correct one.*

What loss does not yet explain

This chapter can say, precisely, how much is lost by a given quotient and exactly where. It has nothing to say about why a particular coarser quotient might be chosen anyway, in full knowledge of the loss it incurs. A designer might prefer a one-bit-lossy traffic light if the saved bit of memory matters more than the distinction it would have preserved; a librarian might prefer the cheap two-category catalog if most patrons never need the finer one. Measuring loss and deciding whether a loss is worth paying for are different questions, and this chapter has only ever answered the first of them.

Chapter 4

Continuation Operators

The first three chapters of this book have all, in their own way, been about nouns: a shape a system of contractions converges to, a partition of histories that share a future, a number measuring what a partition throws away. None of them named the verb. A contraction does something to a shape. A design decision, clearing a bit on arrival at a red light rather than while sitting at one, does something to a history. This chapter is about the doing: what counts as a legitimate move in this framework, and what survives when two different moves are composed into one.

The two kinds of move already used in this book look, on the surface, like opposites. Chapter 1's contractions built detail up, each application adding a smaller, self-similar copy of the whole. Chapters 2 and 3's quotients tore detail down, collapsing histories that had become, for every purpose this book has formalized, the same. It would be a missed opportunity to leave them looking like two unrelated tools. They are both instances of one idea, and this chapter's job is to say what that idea is.

Operators and their composition

Definition 4.1 (Continuation operator). *A continuation operator on X is a Lipschitz map $T : \mathcal{K}(X) \rightarrow \mathcal{K}(X)$: there is a constant $L \geq 0$ with $d_H(T(A), T(B)) \leq L d_H(A, B)$ for all $A, B \in \mathcal{K}(X)$. Write $L(T)$ for the smallest such constant.*

Proposition 4.1 (Closure under composition). *If T_1, T_2 are continuation operators, so is $T_2 \circ T_1$, with $L(T_2 \circ T_1) \leq L(T_2) L(T_1)$.*

Proof. $d_H(T_2(T_1(A)), T_2(T_1(B))) \leq L(T_2) d_H(T_1(A), T_1(B)) \leq L(T_2) L(T_1) d_H(A, B)$. □

Remark 4.1. *This proof is one line because it has to be; the definition was chosen precisely so that composition would cost nothing further to verify. What makes the definition worth having is not this proposition but what it is a common home for. The contractions of Chapter 1 are continuation operators with $L(T) < 1$. The map $\Psi(A) := \bigcup_{z \in A} \Phi(z)$ used throughout Chapters 2 and 3 to advance a reachable set by one step is a continuation operator too, generally with $L(\Psi) \geq 1$: it need not contract, but it is still Lipschitz whenever the underlying one-step map Φ is well behaved. Both belong to the same category. What distinguishes a generative operator from a reductive one is not the definition above; it is what happens under repeated composition, contraction toward a single attractor in one case, and, as the next section makes precise, collapse of distinctions in the other.*

Operators on histories, and what they respect

The operators above act on shapes directly. A different and equally natural kind of operator acts on histories: appending one more real transition, editing a codebase, applying a repair. The interesting question for an operator of this kind is not its Lipschitz constant but whether it interacts predictably with $\sim_\Gamma$.

Definition 4.2 (Γ -compatible operator). $T : H \rightarrow H$ is Γ -compatible if there is a map $\hat{T} : \mathcal{K}(X) \rightarrow \mathcal{K}(X)$ with $\Gamma(T(h)) = \hat{T}(\Gamma(h))$ for every $h \in H$: the continuation space of the transformed history depends on the original history only through its continuation space, not otherwise.

Proposition 4.2 (Γ -compatible operators descend to the quotient). If T is Γ -compatible, then $h_1 \sim_\Gamma h_2 \implies T(h_1) \sim_\Gamma T(h_2)$: T induces a well-defined map $\bar{T} : H/\sim_\Gamma \rightarrow H/\sim_\Gamma$.

Proof. This is Proposition 4 of the previous chapter (quotient invariance), applied to $F := \hat{T} \circ \Gamma$ in place of Γ itself: $h_1 \sim_\Gamma h_2$ gives $\Gamma(h_1) = \Gamma(h_2)$, hence $\Gamma(T(h_1)) = \hat{T}(\Gamma(h_1)) = \hat{T}(\Gamma(h_2)) = \Gamma(T(h_2))$, so $T(h_1) \sim_\Gamma T(h_2)$. $\square$

Remark 4.2. This is the sense in which Γ -compatibility is the right condition to ask for: it says an operator cannot secretly depend on information that continuation equivalence has already declared irrelevant. An operator failing this condition would be able to treat two continuation-equivalent histories differently for reasons that have nothing to do with what either history can still become, which is exactly the kind of dependence Chapter 5 of Continuation Geometry worried about under the name of history-dependence in the first place.

Example 4.1 (The traffic light's step operators). For $e \in \{0, 1\}$, let $T_e : H \rightarrow H$ append the transition $\Phi(\cdot, e)$ to a history. Whether T_e is Γ -compatible (at horizon $T = 1$) needs checking only on the one nontrivial $\sim_\Gamma$ -class found in Chapter 2, $\{(3, 0), (3, 1)\}$, since every other class is a singleton and compatibility there is automatic. Direct computation gives $T_0(3, 0) = T_0(3, 1) = (0, 0)$ and $T_1(3, 0) = T_1(3, 1) = (0, 0)$: both step operators send the two equivalent histories to the identical resulting state, not merely to states with equal R_1 , so compatibility holds here in the strongest possible form.

Remark 4.3. This is a narrow check, confined to the single class where there was anything to check, and it should not be read as a general fact about T_0 or T_1 on this system, only as a confirmation that neither step operator introduces a new dependence on which of $(3, 0)$ or $(3, 1)$ a history happened to pass through, at exactly the point where such a dependence would have mattered.

The family Γ_T and the operator Ψ

A different relationship, easy to conflate with Γ -compatibility but distinct from it, has been used without a name since Chapter 2. It concerns not an operator on histories but the relationship between the finite-horizon continuation spaces at consecutive horizons for one and the same history.

Remark 4.4 (Ψ relates horizons, not histories). By construction, $R_{T+1}(x) = \bigcup_{z \in R_T(x)} \Phi(z) = \Psi(R_T(x))$, i.e. $\Gamma_{T+1} = \Psi \circ \Gamma_T$ as an identity between functions $H \rightarrow \mathcal{K}(X)$, where $\Gamma_T(h) := R_T(x(h))$. This is not an instance of Proposition 4.2; no operator on H is involved, only the fixed

continuation operator Ψ acting on the shape side, once for each increment of the horizon. It is this identity, not any property of an operator on histories, that did the real work in Lemma 2 of the previous chapter (one-step agreement propagates) and in the eventual-periodicity argument of Chapter 3: both are, in retrospect, statements about the orbit of a single reachable set under repeated application of the one continuation operator Ψ , dressed in the language of horizons rather than the language of operators used here.

Two families, one category

Corollary 4.1. *Both the contractions of Chapter 1 and the operator Ψ of Chapters 2 and 3 are continuation operators in the sense of this chapter's first definition; they differ only in whether their Lipschitz constant is below 1.*

Remark 4.5. *This is worth stating plainly because it says the split this book opened with, generative operators building attractors versus reductive operators collapsing distinctions, is not a split in what kind of mathematical object is involved. It is a split in one number. A continuation operator with $L(T) < 1$ pulls every pair of shapes closer together under repeated application, which is exactly Chapter 1's mechanism for generating a unique self-similar attractor. A continuation operator with $L(T) \geq 1$ need not pull anything together at all, and Ψ in particular does something else entirely on a finite state space: not converge to a single shape, but cycle, as the traffic light's own eventual periodicity already showed. Generation and collapse are not two theories. They are two regions of the same parameter.*

What operators do not yet explain

This chapter has said what a legitimate move is and shown that both of this book's central constructions are instances of it, but it has not asked which finite collections of operators are rich enough to produce every continuation geometry worth having, only that individual operators compose predictably once they exist. That question, closer to the universal-generation intuition this project began with, needs the vocabulary built here as a prerequisite, and is not attempted in this chapter.

Chapter 5

Toward Universal Generation

A handful of letters generates every word in a dictionary that has not been written yet. A handful of chess rules generates every game that will ever be played under them, including games no one has conceived of. A small, fixed vocabulary reaching an unbounded space of outcomes is one of the oldest tricks available to any generative system, and it is worth asking, now that this book has a settled vocabulary for what a continuation operator is, whether continuation geometries have the same property. Can a small, fixed family of operators, chosen once and never enlarged, get arbitrarily close to any admissible continuation geometry at all?

This chapter does not answer that question. It draws its boundary precisely enough to say exactly what would need to be true for the answer to be yes, proves the one direction that is actually within reach, and is honest about the direction that is not. The gap between those two is real mathematics, not a gap in effort, and closing it is left where it belongs: open.

The conjecture, stated precisely

Conjecture 5.1 (Universal continuation generation). *There exists a finite family of contractions $T_1, \dots, T_n$ on X , fixed once and independent of any particular target, such that for every nonempty compact $A \subseteq X$ and every $\varepsilon > 0$, some composition $T_{i_k} \circ \dots \circ T_{i_1}$ applied to some starting set satisfies $d_H(A, T_{i_k} \circ \dots \circ T_{i_1}(\cdot)) < \varepsilon$.*

Remark 5.1. *This is stated as a conjecture and not as a theorem because it is one; nothing in this book has proved it, and nothing in this chapter will. What can be made precise is why it is hard, and that requires first showing what is not hard.*

The Collage Theorem

Theorem 5.1 (Collage Theorem). *Let $T_1, \dots, T_n$ be contractions with $\mathcal{H}(A) = \bigcup_i T_i(A)$ and maximum contraction constant $c < 1$, and let A^* be the attractor of $\mathcal{H}$. Then for any nonempty compact A ,*

$$d_H(A, A^*) \leq \frac{d_H(A, \mathcal{H}(A))}{1 - c}.$$

Proof. By the triangle inequality and $A^* = \mathcal{H}(A^*)$,

$$d_H(A, A^*) \leq d_H(A, \mathcal{H}(A)) + d_H(\mathcal{H}(A), \mathcal{H}(A^*)).$$

By Proposition 3 of Chapter 1 ($\mathcal{H}$ is a contraction on shapes), $d_H(\mathcal{H}(A), \mathcal{H}(A^*)) \leq c d_H(A, A^*)$. Substituting,

$$d_H(A, A^*) \leq d_H(A, \mathcal{H}(A)) + c d_H(A, A^*),$$

and rearranging, since $c < 1$, gives $(1 - c) d_H(A, A^*) \leq d_H(A, \mathcal{H}(A))$, which is the claim. $\square$

Remark 5.2. *Nothing here was invented for this chapter; the theorem is a direct three-line consequence of facts Chapter 1 already established, and it is worth noting how much it buys for how little it costs. It converts a question about the attractor, an object defined only as a limit, so generally hard to reason about directly, into a question about a single application of $\mathcal{H}$ to a set already in hand. Finding an operator family whose attractor is close to a target A reduces to finding one whose one-step collage $\mathcal{H}(A)$ is already close to A , an ordinary geometric fitting problem rather than a limiting one.*

Existence of an approximating family

Theorem 5.2 (Every compact target is approximable). *For every nonempty compact $A \subseteq X = \mathbb{R}^n$ and every $\varepsilon > 0$, there exists a finite family of contractions (depending on A and ε) whose attractor A^* satisfies $d_H(A, A^*) < \varepsilon$.*

Proof sketch. Fix $c \in (0, 1)$ and let $\{x_1, \dots, x_m\}$ be a finite δ -net of A for δ small relative to $\varepsilon(1 - c)$, which exists since A is compact, hence totally bounded. Fix a bounded universe $U \supseteq A$ and choose contractions T_i , each with ratio c and fixed point x_i , scaled so that $T_i(U)$ lies within a small neighborhood of x_i (achievable by taking the contraction ratio applied to U small enough relative to δ , which can be arranged independently for each i while keeping the overall maximum ratio bounded away from 1). The resulting collage $\mathcal{H}(A) = \bigcup_i T_i(A)$ then lies within a bounded multiple of δ of the net $\{x_1, \dots, x_m\}$, which is itself within δ of A , giving $d_H(A, \mathcal{H}(A))$ small relative to $\varepsilon(1 - c)$. Theorem 5.1 then gives $d_H(A, A^*) < \varepsilon$. This is the standard construction underlying fractal image compression, given in full quantitative detail in [2], and is not reproduced with complete constants here. $\square$

Remark 5.3. *This is real progress, and it is worth being precise about exactly how much. It says the space of Hutchinson attractors is dense, in Hausdorff distance, in the space of nonempty compact subsets of X : nothing worth generating is permanently out of reach. What it does not say, and what Conjecture 5.1 actually asks for, is that a single family, fixed before any target is known, suffices for every target. The construction above produces a different family for each new target and each new tolerance; as $\varepsilon \rightarrow 0$ or as A grows more intricate, the number of generators m used in the proof generally grows without bound. The existence theorem answers “can this be reached at all,” which Theorem 5.2 settles completely. The conjecture asks “can it be reached with tools chosen in advance,” which is a categorically stronger demand, and this chapter has not moved that needle.*

What kind of gap this is

The distinction is not a technicality, and it has a name in neighboring fields worth being explicit about rather than gestured at. The Stone–Weierstrass theorem guarantees that polynomials are dense in the continuous functions on a compact interval: a fixed class of building blocks, closed under the operations that generate it, dense in everything worth approximating. A universal approximation theorem for neural networks says something structurally similar:

networks built from a fixed activation function, with enough width, are dense in the continuous functions. Both are theorems about a fixed generator class being rich enough for every target, which is exactly the shape of Conjecture 5.1 and exactly what Theorem 5.2 does not provide, since its generator family is chosen after the target is known.

Remark 5.4. *Whether an analogue of Stone–Weierstrass or universal approximation holds for continuation geometries specifically is not something this chapter’s methods can settle, and it should not be assumed to transfer just because the surface shape of the claim resembles those theorems. Both classical results rely on structure, algebraic closure under multiplication for Stone–Weierstrass, compositional depth for neural networks, that has no established analogue here. Building that structure, or showing it cannot exist, is what would actually resolve Conjecture 5.1, and it remains the open question this book set out to sharpen rather than close.*

What Remains

This book is unfinished. What follows is a roadmap, not a draft, so that continuation from here does not require reconstructing the plan from scratch.

Chapter 6 (planned): State Insufficiency, Made Explicit

The sleep/coma motivation used informally since the project's earliest notes has never been built into an explicit toy system. The chapter owed is: construct an actual (X, F, K) with two states x_1, x_2 (or two histories h_1, h_2 landing on the same observable state) such that $\Gamma(x_1) \neq \Gamma(x_2)$, proved, not narrated. This is largely already implicit in *Continuation Geometry*, Chapter 5 (the sufficient-statistic material) and in this book's traffic-light examples, so the honest options are: (a) a short chapter distilling the sharpest existing instance into a clean standalone theorem, framed explicitly as "state insufficiency," or (b) a note explaining why a separate chapter is not needed, since Chapter 5 of the first book already proves the general case. Decide which before writing anything new.

Chapter 7 (planned): Semantic Manifolds

The most speculative and least-earned material still on the table. The goal, meaning as a continuation-equivalence class rather than a point in a fixed space of meanings, requires a specified equivalence before anything else: candidates include $d_H(\Gamma(h_1), \Gamma(h_2)) = 0$ (reusing $\sim_\Gamma$ directly, in which case this chapter may collapse into a corollary of Chapter 2 rather than needing new machinery) or some notion of isometry or homeomorphism between continuation spaces (a genuinely stronger and harder-to-work-with relation, not yet justified over the simpler one). Resist writing a Theorem with an undefined $\cong$ in it. Settle the equivalence first; if it turns out to just be $\sim_\Gamma$, say so plainly rather than introducing new notation to relabel an old idea.

Chapter 8 (planned): Closing

A synthesis chapter in the register of *Continuation Geometry*'s Chapter 11: what was established (Collage Theorem, quotient invariance, the operator/attractor unification of Chapter 4, the loss measure of Chapter 3), what was not (the universal generation conjecture, semantic manifolds if Chapter 7 does not resolve it), and an honest final assessment of whether "generative" turned out to be the right word for what this book actually built.

Loose threads not yet assigned to a chapter

Two things surfaced during drafting and were deliberately not chased: whether operators named Collapse or Refuse elsewhere in this corpus are literally contractions in the sense of Chapter 1 (an empirical question about those specific constructions, not a mathematical one this book can settle from the armchair); and whether a coordination- geometry extension of Chapter 4's operator vocabulary, multiple agents' operators acting jointly on a shared continuation space, is worth its own chapter or belongs in a third volume alongside *Continuation Geometry's* own unfinished coordination-geometry business. Neither should be started without first deciding whether it belongs in this book at all.

Bibliography

- [1] K. Falconer, *Fractal Geometry: Mathematical Foundations and Applications*, 2nd ed., Wiley, 2003.
- [2] M. Barnsley, *Fractals Everywhere*, Academic Press, 1988.
- [3] J. E. Hutchinson, "Fractals and self-similarity," *Indiana University Mathematics Journal*, 30 (1981).